


2026

The year of the AI agent

How AI agents are reshaping governance, risk, and digital trust



The challenge of identity security is rapidly growing in both scale and complexity

Organizations must provide safe and efficient access to an ever-changing tapestry of global employees and contractors. Now, they must also secure a torrent of non-human identities—most notably, autonomous AI agents—working behind the scenes around the clock. Securing access for these identities at scale without slowing down business operations is a massive challenge.

While inefficient human onboarding processes and standing permissions have always opened doors to risk, governing these non-human actors introduces unprecedented challenges. The uncertainty in the market stems from the fact that this is entirely new territory. To truly grasp the scope of the problem, we must define AI agents into their rightful categories, as enterprise risk fundamentally correlates with an agent's level of autonomy and how it is deployed. For example:

- **Assisted & packaged agents (Lower risk):** These act as everyday copilots. They often function as “Client” agents, leveraging existing enterprise credentials to access internal systems. Because a human remains strictly in the loop to drive execution, the security footprint is relatively contained.
- **Low-code & custom-coded agents (Variable to high risk):** These empower businesses to automate complex workflows, but the risk profile dramatically spikes depending on their data access and whether they hold broad, standing permissions. This risk is heavily compounded by the rise of “Shadow AI”—where employees quietly bypass IT protocols to utilize external LLM tools or download unauthorized local agents directly to their machines, creating massive visibility gaps.
- **Fully autonomous & multi-agent systems (Highest risk):** Representing the true “agentic workforce,” these sophisticated entities operate independently in the cloud or across third-party partner ecosystems. With minimal-to-no human oversight, these interconnected agents can self-plan, collaborate, and adapt. They introduce the most profound organizational risks, including cascading machine-speed security breaches, untraceable compliance violations, and unpredictable emergent behaviors.

Managing this diverse spectrum of AI agents requires novel approaches, and frankly, many organizations are unaware of the underlying security issues because there are very few good examples being set. As a result, the traditional access models that once worked for human-only environments are now obsolete and may not keep up with the vast mix of human and machine identities moving across today's enterprise environments.

To understand how organizations are preparing to secure this new landscape in 2026—a year rapidly defining itself as the year of the AI agent—we surveyed 100 security and identity leaders from Fortune 500 and Global 2000 companies.

The adoption of AI agents is moving quickly, but if you're still figuring out how to securely govern this autonomous workforce, you have not fallen too far behind just yet. While the vast majority of organizations are rushing to deploy agents, our research shows that governance policies are still in development at 57% of those same companies. However, the market is sitting at the precipice of a major inflection point. As these non-human identities multiply and gain access to highly sensitive data, now is the time to get serious about this shift before the market leaves you behind. Any organization utilizing AI without proper governance is already vulnerable; as the volume of these agents grows, the probability that a machine-speed breach or compliance failure actually happens grows higher and higher.

Let's walk through some of the most pressing agent-related security challenges and market takeaways that can help your organization as it figures out how to approach this new frontier. We'll explore the current realities of AI agent deployment, the dangerous visibility gaps preventing secure innovation, and how an identity-centric framework—underpinned by robust data access security—provides a secure, compliant path forward for both human and non-human identities.

Executive summary

A fundamental shift is taking root across the enterprise as organizations move beyond conversational AI and begin actively deploying a new breed of action-oriented AI agents. While still in its early stages, this transition introduces unprecedented security challenges, identity blind spots, and a complex web of governance and compliance demands that many organizations are unprepared to meet.

With 91% of organizations already deploying or evaluating AI agents, this is no longer a future concern—it is an immediate reality. This urgency is driven by two colliding forces: first, employees are rapidly creating and using unsanctioned “shadow AI” agents on their own; and second, enterprise SaaS vendors are aggressively “agent-ifying” their existing platforms faster than IT departments can react, evaluate, or disable them.

91%

of organizations are
deploying or evaluating
AI agents

To help organizations navigate this shift, this report explores these urgent challenges using recent industry survey data and introduces an identity-centric framework designed to solve them by providing:

- Discovery and unified visibility to see agents and help satisfy audit requirements.
- Lifecycle management and clear ownership to drive accountability for agent behavior from creation through retirement.
- Real-time governance and audit to control agent access and help prove compliance.
- Proactive protection and response to stop threats before they escalate into reportable incidents.

Without this foundation in place, the rush to innovate is currently creating four core problems for enterprises:

- A visibility gap due to unmanaged shadow AI and a lack of a unified agent inventory, making it difficult to demonstrate comprehensive control to auditors.
- A critical governance deficit, as agents are granted excessive access without clear ownership or adherence to the principle of least privilege, which may violate core tenets of SOC 2 and GDPR.
- A SecOps vulnerability, characterized by the inability to proactively protect against and respond to AI agent threats that operate at machine speed, rendering manual oversight obsolete.
- A compliance breakdown, where traditional, point-in-time compliance checks fail to keep pace with the dynamic, autonomous actions of these new machine identities.

Furthermore, global regulatory frameworks are rapidly evolving to ensure these technologies are deployed safely. **As noted by Gartner**, with AI regulation projected to cover 50% of the world's economies by 2027, organizations need to be prepared. The reality is that organizations lacking agile, automated, and identity-centric governance programs already face significant compliance risks, financial penalties, and stalled innovation today. Because regulations apply equally whether an actor is human or machine, the rise of AI agents simply represents a massive expansion of the non-human identities (NHIs) that must now be brought under these mandates.

By anchoring AI security in identity, organizations can build the agile, compliant, and secure foundation required to innovate with confidence in the agentic era.


The agentic era is here: A new workforce, a new identity crisis

Generative AI was just the prologue; the real disruption lies in autonomous action. We have crossed the threshold from AI that merely converses to AI that executes—systems that independently log in, authenticate, manipulate data, and navigate critical enterprise infrastructure without human intervention. This shift defines the “agentic era,” a fundamental operational pivot that promises massive productivity gains. However, realizing the vision of a fully functioning “agentic workforce”—where complex multi-agent systems collaborate and execute enterprise-wide operations—requires proving these systems are secure through rigorous behavior testing, lifecycle management, and continuous governance cycles.

Currently, the race to deploy these agents is breaking traditional, human-centric security and compliance models. To understand why, consider the mechanics of standard identity tools: traditional Multi-Factor Authentication (MFA) relies on human interaction, like a biometric scan or an SMS code, which a background autonomous agent cannot fulfill. Similarly, point-in-time access reviews, which typically happen quarterly, are practically useless against machine identities that can autonomously spin up, request elevated privileges, execute a task, and spin down in a matter of seconds. When an agent goes rogue or makes an error at machine speed, legacy systems may lack the framework to immediately assign human accountability or revoke access.

Despite these critical security risks, the business drivers pushing for AI agent adoption are incredibly compelling. Organizations are highly motivated by the automation of repetitive tasks (cited by 85% of respondents as a top-3 business outcome), faster or better decision-making (67%), and cost reduction (62%).

What is driving the rush to AI agents?



Percentage of respondents who cited this as a top-3 business outcome

Fueled by these business drivers, adoption is occurring at an unprecedented scale. Ninety-one percent of respondents are either already using AI agents in production or actively evaluating them, with 47% having agents operate in production environments. Organizations draw from a mix of channels to source these agents, including packaged assistants in enterprise software (69%), in-house development (56%), and third-party vendors (54%). As these non-human identities quickly outpace the number of human employees, the rush to deploy them creates a fundamental conflict between rapid innovation and enterprise security.

The urgent problem: Why AI agents create a massive security blind spot


AI agents introduce a new class of identity risk that most organizations are not prepared for, creating pervasive security concerns.

Top 3 concerns



Data security and privacy is a top concern, with 69% of respondents ranking a data breach as a top-3 security priority. The issue is closely followed by unauthorized access to sensitive systems (55%) and, critically, compliance or regulatory violations (47%).

These are not abstract fears. According to survey takers, the majority of AI agents (73%) access sensitive customer data. Nearly half also access HR data (48%) and financial records (43%). With only 3% of organizations reporting that their agents access no sensitive data, the attack surface has expanded dramatically.



97%

of all AI agents have access to sensitive data


This risk is magnified by a profound lack of visibility and control. Only 30% of organizations have complete, real-time visibility into their agentic workforce, leaving the rest with dangerous blind spots. “Shadow AI” is a major part of this problem, as 44% of respondents lack formal processes for detecting unsanctioned AI tools.

Only 30%

have complete, real-time visibility into all AI agents

Visible AI usage

Hidden risks



Beyond data exfiltration and compliance fines, the actual operationalization of these agents is frequently overlooked, leading to devastating system impacts. When non-human identities are deployed without strict lifecycle controls, clear ownership, or behavioral guardrails, their ability to execute tasks at machine speed can quickly overwhelm APIs, exhaust resources, and trigger critical system failures.

Compounding this visibility gap, access controls are woefully inadequate. While most organizations are attempting to enforce the principle of least privilege, our research shows the vast majority are failing to do so completely. A staggering 69% of organizations have only partially implemented least-privilege access for their AI agents, leaving them dangerously exposed. The result is a startlingly small number of organizations that have gotten it right: a mere 11% have fully implemented least-privilege access for AI agents.

11%

Only 11% of organizations have fully implemented least-privilege access for AI agents

This lack of foundational controls, combined with the enterprise rush to adopt autonomous technologies, has created a massive Deployment–Confidence Gap. As noted earlier, while an overwhelming 91% of organizations are actively deploying or evaluating AI agents, this rapid adoption is vastly outpacing their ability to secure it. When asked to rate their “ability to manage AI agent security risks,” only a small fraction—just 21%—report being “very” or “extremely” confident. The vast majority (79%) are moving forward with only moderate, low, or zero confidence in their risk management capabilities.

This deficit extends beyond visibility into a deeper, structural governance crisis. The very frameworks for managing these new digital workers are either incomplete or fundamentally broken. For instance, our data shows that the concept of ownership—the cornerstone of accountability—is dangerously diffused. A staggering 65% of organizations assign AI agent ownership to entire teams rather than a specific individual. While collaborative ownership structures can work, they require clearly defined roles and responsibilities; without them, proving compliant oversight to auditors becomes incredibly difficult and forensic investigations can stall when an agent performs an unauthorized action. Another 17% have no clear ownership structure at all, leaving high-privilege agents to operate as completely orphaned identities. Every AI agent, regardless of whether it is managed by a person or a team, requires a structured and well-defined ownership model tied to its lifecycle to ensure continuous accountability.

When drilling down specifically into AI agents over the past 12 months, 18% of organizations have experienced direct adverse outcomes, including the sharing of sensitive information (44%), unauthorized system access (39%), and the exposure of credentials (39%). With governance policies still in development at 57% of organizations and ownership structures often unclear, the foundation for accountable, auditable AI usage is alarmingly weak.

This is not just a security failure; it is a business inhibitor. The inability to ensure and prove compliance for AI agents stalls innovation, increases the risk of regulatory fines, and erodes customer trust. Without a new framework for identity, the promise of the agentic era will remain unrealized, crippled by the risks it creates.



The SailPoint Agentic Fabric: Your definitive foundation for secure AI innovation

A secure AI strategy is difficult without a unified identity security strategy. Compliance cannot be an afterthought; it must be built into the fabric of your AI infrastructure. The SailPoint Agentic Fabric is an identity-centric product that extends enterprise-grade identity security to the new era of enterprise AI. It anchors everything in identity as the “control plane” for managing access, ensuring that governance keeps pace with innovation.

Built on the proven SailPoint Platform, the Agentic Fabric is guided by six foundational principles for modern agentic security:

- **Immutable birth-right ownership:** Every agent must have a designated human owner to ensure accountability for its actions, creating a clear chain of custody for auditors.
- **Real-time agent ledger:** A complete, unchangeable trace of all agent activities provides an audit trail that can help your organization’s compliance goals and forensic analysis.
- **Continuous risk evaluation:** Access is dynamic; interactions are evaluated for risk in real-time to adjust privileges and prevent standing access.
- **Real-time authorization:** A zero-trust approach grants the minimum privilege at the exact moment needed, enforcing the principle of least privilege by default.
- **Agent moral code:** A final safeguard of core ethical and operational rules provides a policy-based backstop to guide agent actions when other controls fail.

These principles are delivered through three core pillars designed to address the challenges of the agentic era.

The three pillars of agentic security



Pillar 1: Discovery and unified visibility: Uncovering your agentic workforce

The foundational challenge is clear: “You cannot secure what you cannot see.” The lack of a centralized agent inventory, rated as a significant challenge by 46% of respondents, makes comprehensive compliance challenging.

The SailPoint Agentic Fabric is designed to be a single source of truth, discovering agents and mapping its complex relationships to provide a complete and auditable inventory. It leverages a vast connector library for automated, continuous discovery across environments—from agentic platforms and SaaS applications to endpoints—to help detect shadow AI where 44% of organizations currently have no formal process. This moves beyond a simple list to create a unified identity graph, linking agents to their human owners, entitlements, and the sensitive data they can access. This data access path visibility is critical for mitigating exfiltration risks and demonstrating to regulators that you have control over how sensitive data is being used.

Pillar 2: Real-time governance and audit: Bringing control to autonomous agents

Agents with excessive privilege are a critical compliance vulnerability. With only 11% of organizations fully implementing least privilege and ownership often diffused across teams, the potential for misuse is high.

The SailPoint Agentic Fabric helps this by treating agents as first-class identities, extending enterprise-grade governance to them. It enforces clear human ownership, ensuring there is always an accountable individual for an agent’s actions. It applies comprehensive lifecycle management, including automated provisioning and de-provisioning of fine-grained entitlements, to prevent orphaned, high-privilege agents. Most importantly, it enforces least-privilege and just-in-time access, an approach desired by 33% of leaders. By providing an audit-ready trail for agent activity, including robust certification campaigns, it directly addresses the 41% of respondents who list audit and compliance reporting as a top-three desired capability.

How Agentic Fabric supports 2026 regulatory mandates

Regulatory framework	Core mandate	How the SailPoint Agentic Fabric helps your alignment goals
EU AI Act <i>(High-Risk AI Systems)</i>	Article 11 & 12: Technical Documentation & Record-Keeping Requires providers of high-risk AI systems to prepare comprehensive technical documentation before deployment and design systems with continuous, automatic logging of events over the system's lifecycle to ensure traceability of operations.	Pillars 1 & 2: The real-time agent ledger provides an audit-ready cryptographic trace of agent actions, credentials used, and data access paths.
EU AI Act <i>(Human Oversight)</i>	Article 14: Human Oversight Mandates that AI systems must be designed so they can be effectively overseen by natural humans to prevent or minimize risks.	Pillar 2: Enforces immutable birth-right ownership, ensuring every active agent is tied to a responsible human owner with automated access certification campaigns.
SOC 2 <i>(Access Control)</i>	CC6.3: Least Privilege & Zero Trust Requires that access to systems and data is restricted to authorized users and processes based on job roles and business needs.	Pillars 2 & 3: Enforces just-in-time real-time authorization, replacing broad permanent permissions with dynamic, micro-segmented agent access.
SEC Cybersecurity Rules <i>(Incident Disclosure)</i>	Material Cyber Incident Disclosure (4-Day Rule) Requires public companies to disclose material cybersecurity incidents within four business days of determining they are material.	Pillar 3: Proactive protection & response detects anomalous agent behavior quickly, immediately triggering automated revoking of privileges to help contain the threat.

Pillar 3: Proactive protection and response: Defending against threats quickly

A compromised agent operates at machine speed, rendering manual responses and periodic compliance checks inadequate. With only 34% of organizations having good visibility into real-time agent behavior, most are flying blind.

The SailPoint Agentic Fabric delivers proactive protection by continuously monitoring agent behavior for anomalies and policy violations. This real-time monitoring, the top desired capability for 53% of respondents, is essential for a modern compliance posture. An AI-powered risk engine calculates dynamic risk scores for both agents and humans, triggering automated, multi-vector responses when threats are detected. This can include revoking entitlements, disabling agents, and transmitting risk signals to your broader security infrastructure. By sharing **governance-grade identity context with SOC teams**, it transforms your security and compliance posture from reactive to proactive.

Conclusion: Innovate confidently with SailPoint Agentic Fabric

The rise of AI agents is inevitable, and with it, the critical need for a robust identity security strategy that is deeply integrated with compliance objectives. The survey data clearly demonstrates the pervasive risks, governance gaps, and urgent need for better solutions in the agentic era.

The fabric transforms identity security from a point of vulnerability into a powerful security and compliance advantage. By extending adaptive identity security to every AI agent, it enables organizations to harness the full power of the agentic enterprise with confidence. SailPoint's unique identity-centric perspective provides unified governance for identities from a single, enterprise-grade platform, ensuring that security and compliance are enablers of innovation, not barriers to it.


Ready to innovate securely? Request a conversation with SailPoint to discuss your top AI projects and learn how our identity security solutions enable rapid, secure, and compliant AI agent deployment.

About this survey


SailPoint surveyed 100 security and identity practitioners at Fortune 500 and Global 2000 organizations to better understand how enterprises are managing privileged access today, where they face the greatest governance and visibility gaps, and how AI and automation are affecting privileged access strategy. Here’s the breakdown of our respondents by job title and industry:

N=100 total respondents sourced from the ETR-Community and screened for relevant decision-making authority over identity security solutions. Polling was conducted from April 22 to April 28, 2026.

Job Title



Industry





About SailPoint

SailPoint (Nasdaq: SAIL) is defining the new era of adaptive identity security. In a world where non-human identities now significantly outnumber humans, our AI-powered platform unifies identity, security, and data intelligence to protect today's enterprise from advanced identity-based threats. We deliver the identity solution that spans both the breadth of identities and the depth of context needed to drive real-time access with confidence. Built on principles like zero-standing privilege and contextualized risk, our SailPoint platform transforms identity from a point of vulnerability into a powerful security advantage. Trusted by many of the world's leading organizations, SailPoint secures the enterprise with intelligent, autonomous identity security.

© 2026 SailPoint Technologies, Inc. All rights reserved. SailPoint, the SailPoint logo and all techniques are trademarks or registered trademarks of SailPoint Technologies, Inc. in the U.S. and/or other countries. All other products or services are trademarks of their respective companies.